

Max Planck Institute for Security and Privacy

Cryptographic Engineering in Post-Quantum Cryptography

Vincent Hwang


November 6th, 2025, National Taiwan University, Taipei, Taiwan

About Myself



- ▶ BSc., NTU CSIE, Taiwan (2016.09.01 ~ 2021.06.31).
 - ▶ ~ 2020: graph theory/algorithms, complexity theory (surveying).
 - ▶ 2020 ~ 2021, cryptographic engineering:
 - ▶ Courses Post-Quantum Cryptography.
 - ▶ Lattice-based cryptography (course).
 - ▶ Assembly optimizations (internship).
- ▶ MSc., NTU CSIE, Taiwan (2021.09.01 ~ 2022.06.31).
 - ▶ Cryptographic engineering (assembly, lattices).
- ▶ PhD (on going), MPI-SP, Germany (2023.01.01 ~ Now).
 - ▶ Cryptographic engineering:
 - ▶ Assembly programming for lattices.
 - ▶ Formal verification of assembly programs.
 - ▶ Elliptic-curve discrete logarithm (ongoing research).

Cryptography



$P_a = P_b$. Challenge: solve (d_a, d_b, P_a, P_b) from (G, Q_a, Q_b) .

Cryptographic Engineering (Traditional Way)





Cryptosystem

Cryptographic Engineering (Holistic Way)




Algorithms



Cryptosystem


Cryptographic Engineering (Holistic Way)




Cryptographic Engineering (Holistic Way)




Cryptographic Engineering (Holistic Way)




Cryptographic Engineering (Holistic Way)




Cryptographic Engineering (Holistic Way)



Cryptographic Engineering (Holistic Way)



Cryptographic Engineering (Holistic Way)



Cryptographic Engineering (Life Cycle)



Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?

FFT
Float
NTT
Modular arith.

Which hashes?

SHA-2
SHA-3

Randomness sampling?

High-level descriptions



Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Power side-channels
Performance (hardware)
Area
Latency
Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization

Post-Quantum Cryptography



- ▶ Pre-quantum: integer factorization (IF) and elliptic-curve discrete logarithm (ECDLP).
- ▶ Post-quantum (believed to be secure against quantum computers):
 - ▶ Lattice-based.
 - ▶ NTRU, LWE, M/R-LWE, M-LWR, M-SIS.
 - ▶ Polynomial multiplications, matrix multiplications over \mathbb{Z}_q ($\mathbb{Z}/q\mathbb{Z}$).
 - ▶ Supersingular-isogeny-based.
 - ▶ Multivariate-based.
 - ▶ Code-based.
 - ▶ Hash-based.
 - ▶ Rich algebraic structures except for the hash-based ones.

NIST Post-Quantum Cryptography Standards



Post-Quantum Cryptography Standardization by the National Institute for Standards and Technology (NIST).

- ▶ Selecting cryptosystems: 2016 ~ 2025. Four rounds.
- ▶ FIPS203 (ML-KEM, Kyber), 2024.
 - ▶ Lattice. Module-Learning-With-Error (M-LWE).
- ▶ FIPS204 (ML-DSA, Dilithium), 2024.
 - ▶ Lattice. M-LWE, Module-Short-Integer-Solution (M-SIS).
- ▶ FIPS205 (SLH-DSA, SPHINCS⁺), 2024.
 - ▶ Hash-based.
- ▶ FIPS206 (FN-DSA, Falcon), 2025?
 - ▶ Lattice. SIS over NTRU, fast Fourier sampling,
- ▶ HQC, 2027?
 - ▶ Code-based.

What Am I Doing?



- ▶ Assembly optimization for polynomial multiplications.
 - ▶ Assembly: [Armv7-M](#) (microcontrollers), Armv8-A and AVX2 (smartphones, personal computers).
 - ▶ Lattices: [Dilithium](#), Kyber, NTRU, NTRU Prime, and Saber.
 - ▶ Polynomial rings:
 $\mathbb{Z}_q[x]/\langle x^{256} + 1 \rangle$, $\mathbb{Z}_{2^{13}}[x]/\langle x^{256} + 1 \rangle$, $\mathbb{Z}_{2^k}[x]/\langle x^n - 1 \rangle$, $\mathbb{Z}_q[x]/\langle x^p - x - 1 \rangle$
 - ▶ [Optimizations for modular arithmetic](#).
 - ▶ Optimizations for polynomial transformations.
- ▶ Formal verification of assembly.
 - ▶ [Floating-point in Falcon \(lattice\)](#).
- ▶ Cryptanalysis.
 - ▶ [ECDLP](#) on data-center-level GPUs.

This talk.

Multiplying Polynomials without Powerful Multiplication Instructions

Vincent Hwang, YoungBeom Kim, and Seog Chung Seo

- ▶ Variable-time long multiplication instructions.
- ▶ Constant-time modular arithmetic.
- ▶ Prime modulus: generalized Barrett multiplication suitable for multi-limb arithmetic.
- ▶ Power-of-two modulus: Nussbaumer.
- ▶ Paper: <https://tches.iacr.org/index.php/TCHES/article/view/11926>.
- ▶ Artifact: https://github.com/vincentvh/PolyMul_Without_PowerfulMul.

Where Are We?



Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?

FFT
Float
NTT
Modular arith.

Which hashes?

SHA-2
SHA-3

Randomness sampling?

High-level descriptions



Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Emulation

Power side-channels
Performance (hardware)
Area
Latency

Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization

What Is Constant-Time in Cryptographic Implementations?



- ▶ Constant-time: ~~constant execution time~~, **secret-independent** execution time.
- ▶ Variable-time: **secret-dependent** execution time.
- ▶ Common computations that leak timing side-channels if input is secret.
 - ▶ Conditional branches.
 - ▶ Table lookup.
 - ▶ Indexing with ~~secret index~~.
 - ▶ Some assembly instructions.
 - ▶ Divisions.
 - ▶ Floating-point arithmetic.
 - ▶ [Long multiplications](#).
 - ▶ Typical resolution: emulate with constant-time instructions.

Modular arithmetic in ML-DSA




- ▶ ML-DSA.
- ▶ Module-lattice-based digital signature.
- ▶ $q = 2^{23} - 2^{13} + 1$. Modular arithmetic in \mathbb{Z}_q .
- ▶ Let `int32_t a b`; we often need one of the following:
 - ▶ A 64-bit product.
 - ▶ `int64_t c = (int64_t)a * b;`
 - ▶ The highest 32 bits of a 64-bit product.
 - ▶ `int64_t c = (int32_t)((int64_t)a * b) >> 32;`
 - ▶ Cortex-M3: 32-bit Armv7-M instruction set architecture (ISA).
 - ▶ 32-bit long multiplications are variable-time → leak secret information.
 - ▶ Multi-limb arithmetic for the long product → significant overhead.

Multiplication Instructions



R is a power of two, 2^{32} on Cortex-M3.




Generalized Barrett Multiplication (2-Limb)




Goal: compute $a \cdot c \equiv ab \pmod{q}$. $\left\lfloor \frac{ab}{q} \right\rfloor \approx \left\lfloor \frac{ab'}{R} \right\rfloor$, $b' = \left\lfloor \frac{bR}{q} \right\rfloor$.

Originally, compute $c = ab - \left\lfloor \frac{ab'}{R} \right\rfloor q$ with $\left\lfloor \frac{ab'}{R} \right\rfloor$ as:



Instead, compute $c = ab - \left\lceil \frac{ab'}{R} \right\rceil q$ with $\left\lceil \frac{ab'}{R} \right\rceil$ as:





- ▶ For a $\delta > 0$, δ -integer approximation $\llbracket \cdot \rrbracket$: $\forall r \in \mathbb{R}, |r - \llbracket r \rrbracket| \leq \delta$.
- ▶ For a b , write $b_l + b_h \sqrt{R} = \left\lfloor \frac{bR}{q} \right\rfloor$. Define $\llbracket \cdot \rrbracket_b$ as

$$\forall r \in \mathbb{R}, \llbracket r \rrbracket_b := \left\lfloor \frac{a_l b_h}{\sqrt{R}} \right\rfloor + \left\lfloor \frac{a_h b_l}{\sqrt{R}} \right\rfloor + a_h b_h$$

where $a_l + a_h \sqrt{R} = \frac{rR}{\left\lfloor \frac{bR}{q} \right\rfloor}$, $b_l + b_h \sqrt{R} = \left\lfloor \frac{bR}{q} \right\rfloor$.

- ▶ Obviously, $|\llbracket r \rrbracket_b - r| < 3$ (see previous slide).
- ▶ $|\llbracket r \rrbracket_b - \lfloor r \rfloor| < 3$ (see paper).
- ▶ $\left| \left(ab - \llbracket \frac{ab}{R} \rrbracket_b q \right) - ab \bmod q \right| \leq 3q$.

Results



- ▶ $1.92 \times$ faster modular multiplication.
 - ▶ \approx as a long multiplication.
 - ▶ One of the operands must be a constant.
 - ▶ An extra precomputed constant.
 - ▶ No assumptions on the modulus.
 - ▶ No assumptions on the constants.
 - ▶ Above conditions are usually met in lattice-based cryptosystems.
- ▶ $1.51 \times$ faster NTT (for poly. mul.) in $\mathbb{Z}_q[x]/\langle x^{256} + 1 \rangle$.

Table: Cycle count on Cortex-M3.

Long	11
Montgomery	23
Barrett	12

Optimize **a modular multiplication** instead of **a long multiplication + a modular reduction**.

Formal Verification of Emulated Floating-Point Arithmetic in Falcon

Vincent Hwang

- ▶ Paper: https://link.springer.com/chapter/10.1007/978-981-97-7737-2_7.
- ▶ Artifact: https://github.com/vincentvh/Float_formal.

Where Are We?



Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?

FFT
Float

NTT

Modular arith.

Which hashes?

SHA-2
SHA-3

Randomness sampling?

High-level descriptions



Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Emulation

Power side-channels
Performance (hardware)
Area
Latency

Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization




- ▶ FIPS206, FN-DSA. Drafting.
- ▶ Lattice-based digital signature scheme.
- ▶ Hash-and-sign.
- ▶ Double-precision floating-point arithmetic in signing.
- ▶ Concerns of floating-point arithmetic.
 - ▶ Not always constant-time.
 - ▶ Double-precision FPU does not even exist on some microcontrollers!
 - ▶ Cortex-M3 (Armv7-M).
 - ▶ Cortex-M4F (Armv7E-M + single-precision).
 - ▶ Emulating with integer and logical arithmetic.

Floating-Point Arithmetic



Double-precision in this talk.



- ▶ $0 < e < 2047$ (normal values):

$$(-1)^s 2^{e-1075} (2^{52} + m).$$

- ▶ Zeros: $e = m = 0$.
- ▶ Other values: $e = 0 \wedge m \neq 0$, $e = 2047$.

Motivations of Formal Verification for Floats



Incorrect zeroization in floating-point multiplication (FP mul.).

- ▶ Emulated floating-point arithmetic (by Falcon submitters) in C and Armv7-M assembly.
- ▶ Zeroizing and rounding in the wrong order.
- ▶ \exists **normal non-zero** floating-point products $\approx \pm 2^{-1074}$ being **zeroized**.
- ▶ 692/2048 (34%) float constants in FFT admit such float products!

Questions.

- ▶ \exists such floats in Falcon? Experimentally no, but no proofs.
- ▶ What are the programs actually doing?



- ▶ \exists such float products in Falcon? Range checking.
 - ▶ No in FFTs of Falcon.
 - ▶ With **input conditions**, $\neg\exists$ floats with incorrect zeroizations.
 - ▶ If FFT inputs are integers drawn from $[-2^{15}, 2^{15}]$, every floats in $[2^{-476}, 2^{27}(2^{52} + 605182448294568)]$ ✓.
 - ▶ 0.3 seconds for an FP addition.
 - ▶ 2.6 seconds for an FP multiplication.
- ▶ What are the programs actually doing? Functional equivalence.
 - ▶ == more readable programs.
 - ▶ Armv7-M assembly \longleftrightarrow CryptoLine \longleftrightarrow Jasmin.
 - ▶ \approx 1 minute for each FP addition \longleftrightarrow .
 - ▶ 5 ~ 60 seconds for each FP multiplication \longleftrightarrow .

More Published Works I



Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?

FFT

Float

NTT

Modular arith.

Which hashes?

SHA-2

SHA-3

Randomness sampling?

High-level descriptions



Armv7-M

Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Power side-channels
Performance (hardware)
Area
Latency
Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization

More Published Works II




Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?
FFT
Float
NTT
Modular arith.

Which hashes?
SHA-2
SHA-3
Randomness sampling?

High-level descriptions



Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Power side-channels
Performance (hardware)
Area
Latency
Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization

More Published Works III



Recommendations

Which parameters?
Module dimension
Polynomial dimension

Coefficient ring

How to compute?

FFT

Float

NTT

Modular arith.


Which hashes?

SHA-2

SHA-3

Randomness sampling?

High-level descriptions



Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Power side-channels
Performance (hardware)

Area

Latency

Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization



Recommendations

Which parameters?
Module dimension
Polynomial dimension
Coefficient ring

How to compute?
FFT
Float
NTT
Modular arith.

Which hashes?
SHA-2
SHA-3
Randomness sampling?

High-level descriptions



Armv8-A Neon, x86 AVX2

Implementations

Timing side-channels
Conditional jumps
Table lookups
Variable-time inst.
Power side-channels
Performance (hardware)
Area
Latency
Performance (software)
Low-level arithmetic
Register pressure
Memory access
Vectorization

Publications



Authors in alphabetic order.

ACCEHHLNSWY20 CHKSSY21 ACCHKY21

Microcontroller opt.

AHKS22 HKS24 BHKPY22

AABEHOPSS25

Formal verification

HLSSTYY22

Survey of opt.

BHKYY21 AHY22 CCHY24 HLY24

H24a H24c

Mid-tier/high-end processor opt.

Ongoing Works



- ▶ Deniable authenticated KEM.
 - ▶ USENIX 2026, rebuttal phase.
 - ▶ Feasibility results in practice (mid-tier/high-end processors).
 - ▶ Coauthors: theory.
 - ▶ My job: implementation.
- ▶ Elliptic-curve discrete logarithm on GPUs.
 - ▶ A random curve over a generic 131-bit prime.
 - ▶ H100 GPU.
 - ▶ Projection: 66,000 H100-days.
 - ▶ Low-level optimizations (inline ptx).
 - ▶ Memory access.



- ▶ Optimizations.
 - ▶ More architectures: Armv9-A and AVX-512.
 - ▶ More cryptosystems: Falcon.
- ▶ Formal verification of Falcon.
 - ▶ Transformation from high-dim. to one-dim. samplers.
 - ▶ Functional correctness.
 - ▶ Range properties.
 - ▶ Security arguments for the samplers.
 - ▶ Rényi divergence.
 - ▶ \exists math. proofs for the one-dimensional. case in literature.
 - ▶ $\neg\exists$ math. proofs for the high-dimensional case.
- ▶ Cryptanalysis on GPUs.
 - ▶ Elliptic-curve method in the general number field sieve for **IF**.
 - ▶ Random curves over random ~ 130 -bit integers.
 - ▶ Lattice sieving for shortest vector problem.
 - ▶ GPU-memory-friendly bucketing.
 - ▶ Inner products.



Thank you for attention