

Max Planck Institute for Security and Privacy

Implementing Polynomial Multiplications for Lattice-Based Cryptography on Microcontrollers

Vincent Hwang

January 6th

Today's Plan



1. Lattice-based cryptography.

- ▶ Well-balanced size and computational efficiency.
- ▶ Popular (well-studied).
- ▶ Computational bottleneck:
 - ▶ Polynomial multiplications.
 - ▶ Cryptographic hash functions.

2. Microcontrollers.

- ▶ Resource-constraint devices.
- ▶ Low-level implementation is more preferred.
- ▶ Optimizing compilers are not very “optimizing.”
- ▶ Susceptible to implementation attacks.

3. Polynomial multiplications.

- 3.1 Modular arithmetic.
- 3.2 Fast homomorphisms (more on memory).

Lattice-Based Cryptography



- ▶ Cryptography relying on **lattice** hard problems.
- ▶ Lattice hard problems:
 - ▶ LWE, R-LWE, M-LWE,
 - ▶ SIS, R-SIS, M-SIS, NTRU,
 - ▶ and more...

Unstructured Learning-With-Errors



1. Generate public A (uniform), secret s, e (small support).
2. Compute public $As + e$.
3. Challenge: find (s, e) from $(A, As + e)$.
 - ▶ A, s, e : large matrices
 - ▶ FrodoKEM (NIST round 3 alternate candidate).
 - ▶ Computational bottleneck: As (dimension $n \times n$).
 - ▶ $n = 640 \sim 1344$.
 - ▶ $O(n^3)$ or $O(n^{\log_2 7})$.
 - ▶ Too slow for microcontrollers.

Structured Learning-With-Errors



1. Generate public A (uniform), secret s, e (small support).
 2. Compute public $As + e$.
 3. Challenge: find (s, e) from $(A, As + e)$.
- ▶ Ring-learning-with-errors (R-LWE).
 - ▶ A, s, e : elements in **a large polynomial ring**.
 - ▶ NewHope (NIST round 2 candidate).
 - ▶ Module-learning-with-errors (M-LWE).
 - ▶ A : small matrix over **a large polynomial ring**.
 - ▶ s, e : small vectors over **a large polynomial ring**.
 - ▶ More flexible on the parameter choices.
 - ▶ ML-KEM, ML-DSA (NIST standards).
 - ▶ We need fast polynomial multiplications!
 - ▶ LWE: $O(n^3)$ or $O(n^{\log_2 7})$.
 - ▶ R-LWE, M-LWE (with small module dimension): $O(n^2)$ or $O(n \log n)$.

Polynomial Multiplication



Compute $a(x)b(x)$ in $R[x]/\langle g(x) \rangle$, $R = \mathbb{Z}_m := \mathbb{Z}/m\mathbb{Z}$ for a positive integer m .

1. Modular arithmetic.

- ▶ How efficient the ring arithmetic in R can be?

2. Fast homomorphisms.

- ▶ Which asymptotically faster (faster than $O(n^2)$) approaches we should choose?
- ▶ Very complicated. Decide at the end.


3. Memory operations.

- ▶ Usually invisible from the math description.

Multiplication Instructions



$R = 2^{16}, 2^{32}$. Signed arithmetic throughout this talk.



- ▶ How to map the modular arithmetic to multiplication instructions?
- ▶ How do the multiplication instructions perform on a particular processor?

Montgomery Multiplication



Fixed modulus q , precision R (power-of-two), precomputed $-q^{-1} \bmod R$.

Given a, b , compute

$$\frac{ab + (-abq^{-1} \bmod R)q}{R} \equiv abR^{-1} \pmod{q}.$$

Division and modulo reduction by R are cheap, free if nicely aligned with the architecture.

- ▶ Result derived from the high parts of the long products.
- ▶ Result is scaled by $R^{-1} \bmod q$.
 - ▶ Additional Montgomery multiplication with $R^2 \bmod q \longrightarrow \equiv ab \pmod{q}$.
 - ▶ Replace b by $bR \bmod q$ (precomputation required).

Cost of Montgomery Multiplication (with Precomputation)



$$\frac{a(bR \bmod q) + (-a(bR \bmod q)q^{-1} \bmod R)q}{R} \equiv ab \pmod{q}$$

1. $a, (bR \bmod q) \mapsto a(bR \bmod q)$: long mul.
2. $a(bR \bmod q) \mapsto -a(bR \bmod q)q^{-1} \bmod R$: low mul. (assuming $\text{mod } R$ is free).
3. $-a(bR \bmod q)q^{-1} \bmod R \mapsto (-a(bR \bmod q)q^{-1} \bmod R)q$: long mul.
4. Add.
5. Divide by R .

Summary: 2 long mul. + 1 low mul.

Aligning the Arithmetic with Low-Level Instructions



- ▶ Armv7-M (32-bit ISA): native 32-bit low (w/ acc.), long mul. (w/ acc.) instructions.
- ▶ Processors:
 - ▶ Cortex-M3.
 - ▶ Low takes 1 cycle, acc. one takes 2 cycles.
 - ▶ Long mul takes $3 \sim 7$ cycles.
 - ▶ Cortex-M4.
 - ▶ DSP extension: high multiplication instructions.
 - ▶ Each multiplication instruction takes 1 cycle.

How many cycles does 32-bit Montgomery multiplication take?

- ▶ 9 \sim 16 cycles on Cortex-M3.
- ▶ 3 cycles on Cortex-M4.

Alternatives?



- ▶ Can we do better than 32-bit Montgomery mul. on Cortex-M3?
 - ▶ Yes, 32-bit Barrett mul is faster.
- ▶ Can we do better than 32-bit Montgomery mul. on Cortex-M4?
 - ▶ No other known approach outperforming it.
 - ▶ Other approaches perform the same, but with more memory for the precomputed values.

Barrett Multiplication



Given a, b , we have

$$ab \bmod q = ab - \left\lfloor \frac{ab}{q} \right\rfloor q$$

by definition.

Barrett mul. approximates $\left\lfloor \frac{ab}{q} \right\rfloor$ up to a small integral perturbation. Various formulations (bR/q precomputed).

- ▶ $\left\lfloor \frac{\frac{a}{R} \lfloor bR/q \rfloor}{R} \right\rfloor$.
- ▶ $\left\lfloor \frac{\frac{a}{R} \lfloor bR/q \rfloor}{R} \right\rfloor$.
- ▶ $\left\lfloor \left[\frac{a}{R} \lfloor bR/q \rfloor \right] \right\rfloor$, advanced rounding technique [1].

Cost of Barrett Multiplication



Typically, $ab - \left\lfloor \frac{a bR/q}{R} \right\rfloor q \in [-\frac{R}{2}, \frac{R}{2})$ so

$$ab - \left\lfloor \frac{a \lfloor bR/q \rfloor}{R} \right\rfloor q = (ab \bmod R) - \left(\left\lfloor \frac{a \lfloor bR/q \rfloor}{R} \right\rfloor q \bmod R \right).$$

1. $a, b \mapsto ab \bmod R$: low mul.
2. $a, \lfloor bR/q \rfloor \mapsto \left\lfloor \frac{a \lfloor bR/q \rfloor}{R} \right\rfloor$: high mul.
3. $\left\lfloor \frac{a \lfloor bR/q \rfloor}{R} \right\rfloor, q \mapsto \left\lfloor \frac{a \lfloor bR/q \rfloor}{R} \right\rfloor q \bmod R$: low mul.
4. Subtract.

Summary: 2 low mul. + 1 high mul.

- ▶ Cortex-M3: 6 ~ 10 cycles.
- ▶ Cortex-M4: 3 cycles.

Comparisons



- ▶ Cortex-M3:
 - ▶ Montgomery: 9 ~ 16 cycles.
 - ▶ (Preferred) Barrett: 6 ~ 10 cycles.
- ▶ Cortex-M4:
 - ▶ (Preferred) Montgomery: 3 cycles, precomputed $bR \bmod q$ (we don't need b anymore).
 - ▶ Barrett: 3 cycles, precomputed $\left\lfloor \frac{bR}{q} \right\rfloor$ (we still need b).
- ▶ Know your instruction set architecture.
- ▶ Know your platform.
- ▶ Constant-time?
 - ▶ Advanced rounding technique on Cortex-M3.
 - ▶ Solution: see [this paper](#):
<https://tches.iacr.org/index.php/TCCHES/article/view/11926/11785>.

Fast Homomorphisms



$\deg \mathbf{g} = n, \deg \mathbf{h}_i < n.$

$$R[x]/\langle \mathbf{g} \rangle \xrightarrow{\text{Fast homomorphism } f.} \prod_i R[x]/\langle \mathbf{h}_i \rangle$$


$$\mathbf{a}(x)\mathbf{b}(x) = f^{-1}(f(\mathbf{a}(x))f(\mathbf{b}(x))) \in R[x]/\langle \mathbf{g} \rangle$$

- ▶ Large polynomial ring \rightarrow several small polynomial rings.
- ▶ Toom– k : $O(n^{\log_k(2k-1)})$.
- ▶ Number-theoretic transforms: $O(n \log_2 n)$.
- ▶ Number-theoretic transforms (with poly. ring extensions): $O(n \log_2 n \log_2 \log_2 n)$.
- ▶ A lot of ways.

Frequently Seen Patterns




Component-wise arithmetic of sub-polynomials.




- ▶ Chop phase is free (for illustration only).
- ▶ Transform phase costs.
 - ▶ Arithmetic cost can be inferred from the math.
 - ▶ Memory cost is usually invisible from the math.

Number-Theoretic Transform



- ▶ Poly. mul. in $R[x]/\langle x^{16} - 1 \rangle \rightarrow$ four poly. mul. of the form $R[x]/\langle x^4 - \alpha \rangle$.
- ▶ $R[x]/\langle x^{16} + 1 \rangle$: 256 mul. + 240 add.
- ▶ Four $R[x]/\langle x^4 - \alpha \rangle$: 64 mul. + 48 add.

Arithmetic Cost



1. $R[x]/\langle x^{16} + 1 \rangle \rightarrow R[x]/\langle x^8 - \omega_4 \rangle \times R[x]/\langle x^8 + \omega_4 \rangle$: 8 mul., 16 add./sub.
 2. $R[x]/\langle x^8 - \omega_4 \rangle \rightarrow R[x]/\langle x^4 - \omega_8 \rangle \times R[x]/\langle x^4 + \omega_8 \rangle$: 4 mul., 8 add./sub.
 3. $R[x]/\langle x^8 + \omega_4 \rangle \rightarrow R[x]/\langle x^4 - \omega_8^5 \rangle \times R[x]/\langle x^4 + \omega_8^5 \rangle$: 4 mul., 8 add./sub.
- ▶ Transformation cost: 16 mul. and 32 add.
 - ▶ Three transformations for poly. mul: 48 mul. and 96 add.
 - ▶ Total: 112 mul. + 144 add. (previously 256 mul. + 240 add.).

Memory and Registers




- ▶ Registers:
 - ▶ Fast memory inside the processor.
 - ▶ # register is ISA-fixed.
 - ▶ Bit-size of register is ISA-fixed.
- ▶ Memory:
 - ▶ Access through memory bus: incurring overhead.
 - ▶ Size of memory depends on what the platform designers want.

How to minimize transfers between memory and registers?


- ▶ Issue as much arithmetic as possible while dependent data are in register.
- ▶ Restructure the computation flow if needed.

Memory Operation Cost (First Try)



Two load-store pairs for each entry. First and Second.

Memory Operation Cost (Optimized): Layer-Merging




One load-store pair for each entry.

Memory Operation Cost (Optimized, Concretely)




22




Memory Operation Cost (Optimized, Concretely)




23



Memory Operation Cost (Optimized, Concretely)



Memory Operation Cost (Optimized, Concretely)



Summary I



Optimize polynomial multiplication on microcontrollers.

1. Pick a modular arithmetic.
2. Pick a fast homomorphism.
 - ▶ A lot more to say. See [this survey](https://cic.iacr.org/p/1/2/1): <https://cic.iacr.org/p/1/2/1>.
3. Optimize the memory operations.

In reality, we jump between each steps.

- ▶ A particular modular arithmetic might be the fastest one for a homomorphism on a certain platform, BUT NOT ON OTHER PLATFORMS.
- ▶ A particular homomorphism might be best applied to the polynomial ring on a certain platform, BUT NOT ON OTHER PLATFORMS.

Summary II



Know your platform.

- ▶ Instruction set architecture (ISA):
 - ▶ **Armv7E-M:** <https://developer.arm.com/documentation/ddi0403/ed>.
 - ▶ **Armv8-M:** <https://developer.arm.com/documentation/ddi0553/latest>.
 - ▶ **Armv7-A:** <https://developer.arm.com/documentation/ddi0406/cb>.
 - ▶ **Armv8-A:** <https://developer.arm.com/documentation/ddi0487/gb/?lang=en>.
 - ▶ **Armv9-A:** <https://developer.arm.com/documentation/ddi0487/latest/>.
- ▶ Processors.
 - ▶ One for each processor.
 - ▶ If unlucky, no document for the target processor.
 - ▶ Benchmark the instructions following the ISA.
 - ▶ If unlucky, undocumented instructions.
 - ▶ AMX in Apple M1, M2, M3: See [here](https://github.com/corsix/amx): <https://github.com/corsix/amx>.
 - ▶ Guess how instructions are encoded.
 - ▶ Benchmark the undocumented instructions.

Summary III



Homomorphisms in ML-KEM and ML-DSA are fixed.

ML-KEM			
Processor	ISA	Modular arith.	Layer-merging
Cortex-M3	Armv7-M	Plantard	3
Cortex-M4	Armv7E-M	Plantard	4
Apple M1	Armv8-A	Barrett	4
Haswell	x86 + AVX2	Montgomery	3

ML-DSA			
Processor	ISA	Modular arith.	Layer-merging
Cortex-M3	Armv7-M	Barrett	3
Cortex-M4	Armv7E-M	Montgomery	4
Apple M1	Armv8-A	Barrett	4
Haswell	x86 + AVX2	Montgomery	3



Thank you for your attention

Slides: https://vincentvvh.github.io/slides/UK_MMU-UTAR_PQC_2026_01_06.pdf